Detecting Fraud on Plastic Card Transactions

Project Report Draft
December 4, 2017

Daniel Casares and Felipe Olivera

Executive Summary (1 page)

– Concise problem statement

– List of major concerns/assumptions (if any)

– Summary of findings

– Recommendations

Problem description and background

The use of plastic cards (i.e. credit and debit cards) as a payment method has grown significantly over past years, unfortunately so has fraud (Bahnson 134). Plastic card fraud is defined as an unauthorized account activity committed by means of the debit and credit facilities of a legitimate account. Some successful fraud tactics observed in the industry are lost and stolen card fraud, counterfeit card fraud, card not present fraud, mail non-receipt card fraud, account takeover fraud and application fraud (Krivko 6070). Based on the latest figures gathered in 2015, card fraud accumulated USD 21.84 Billion worldwide in losses (The Nilson Report 6). When banks lose money due to credit card fraud, the losses partially are passed to customers through higher interest rates, higher membership fees and reduced benefits. Hence, it is both the banks' and cardholders' interest to reduce illegitimate use of credit cards (Maes 2).

In this work, we consider the problem of identifying whether a credit or debit card account has been subject to fraudulent activity, using real-life transaction data from a Latin American credit card processing company. The goal is to construct a supervised learning model that can detect fraud on new (previously unseen) plastic card transactions. Fraud detection is, given a set of credit card transactions, the process of identifying those transactions that are fraudulent. Thus, the transactions are classified as genuine or as fraudulent transactions (Maes 2). Different detection systems that are based on machine learning techniques have been successfully used for this problem, in particular: neural networks, bayesian learning, artificial immune systems, association rules, hybrid models, support vector machines, peer group analysis, decision tree techniques such as ID3, C4.5, and random forest, discriminant analysis, social network analysis and logistic regression (Bahnson 135, Mahmoudi 2510).

Exploratory data analysis

For this project we used one dataset provided by a Latin American card processing company. Before explaining the details of the dataset is important to state that it isn't available in public sources (Kaggle, UCI Machine Learning Repository, etc.), it was obtained by "Universidad Católica del Uruguay" in a collaboration project with a Latin American card processing company. The dataset consists of fraudulent and legitimate transactions made with debit and credit cards between July 2014 and June 2015. The total dataset contains 41,091,288 individual transactions, each one with 13 attributes, including a fraud label indicating whenever a transaction is identified as fraud (Table 1). This label was created internally in the card processing company, and can be regarded as highly

accurate. In the dataset only 12,632 transactions were labeled as fraud, leading to a fraud ratio of 0.031%.

Attribute name	Description
amount	Amount of the transaction in USD
id_issuer	Unique identifier of the bank issuer of the card
id_merchant	Unique identifier of the merchant
datetime	Date and time of the transaction
country_code	Numeric code that identifies the country of the transaction
tokenized_pan	Unique identifier of the credit card (Primary Account Number (PAN))
pos_entry_mode	Numeric code that identifies the transaction entry mode (e.g. Chip and PIN, magnetic strip, etc.)
id_mcc	Identification of the Merchant Category Code (ISO 18245)
is_upscale	Indicates if the card holder is an upscale customer
mcc_group	Merchant Category Code grouping by major type of business
type	„C" for credit cards, „D" for debit cards
is_fraud	1 if the transaction was fraudulent, 0 otherwise

Table 1: All 13 attributes of the dataset

Due to the low proportion of the target class (i.e. frauds) in the given dataset, the class imbalance problem arises. Classification of imbalanced data is difficult because standard classifiers are driven by accuracy, thus the minority class may simply be ignored (Visa 67). Generally all classifiers present some performance loss when the data is unbalanced (Prati 253). Additionally, many imbalanced datasets experience problems related to its intrinsic characteristics, such as lack of density and information. To illustrate these issues, a dataset containing of 5 : 95 minority-majority examples and a dataset of 50 : 950 are compared. Though the imbalance factor is the same as in both datasets in the first case the minority class is poorly represented and suffers more from the lack of information factor than in the second case.

Another difficulty associated with this dataset is the computing power required to preprocess the data and train the different predictive models, due to large number of observations contained in it. Since we haven't enough computing power, we decided to under-sample the dataset. Our approach was to include all the frauds of the original dataset. To do that, we loaded the 41 million records in R, and selected the list of "tokenized_pan" for all the fraudulent transactions ("is_fraud" equals to 1), resulting in 3,841 unique customers. Then, from the whole dataset, 8,000 customers were randomly selected and added to the previously mentioned list. After executing the *unique*

function again over the list, it contains 11,296 unique tokenized PANs. By selecting all the transactions associated with those card numbers, the *new* dataset contains 523,049 transactions and a fraud ratio of 2.33%. This process not only resulted in a smaller dataset, decreasing the computing power needed to work with it, it also decreased the imbalance problem (increasing the minority class proportion more than 75 times). In this *new* dataset, the total financial losses due to fraud are USD 1,876,697 . From plotting the amount of fraudulent and total transactions over time we can see that the proportion of fraudulent transactions varies over time (Figure 1). In Dec 2014 the fraud rate was the highest. However, we have no explanation for this behavior.


Figure 1: Amount of fraudulent and total transactions over time

The boxplot of the log amount per transaction in USD versus non fraudulent and fraudulent transactions shows that the amount of fraudulent transaction is a bit higher and with less extreme values (Figure 2). So, they were neither extreme low nor extreme high. The natural logarithm of amount was taken because amount has a lot of high outlying values making a meaningful interpretation of the boxplot impossible.


Figure 2: Log amount per transaction in USD versus non fraudulent and fraudulent transactions

Analysis plan

Explanation of modeling choice

In scientific literature three basic tried and tested classification algorithms are discussed. These are logistic regression, decision tree and random forest (Whitrow 31-51, Bahnson 134). Apart from these three models we implement another tree model, the *Extreme Gradient Boosting (XGBoost)* that is often a winning model for data science competitions on *Kaggle* (Gordon). For modeling purposes the *new* dataset is split randomly in 70% training and 30% test taking into account the same proportions of class labels in both data sets. The models are implemented in R and trained and validated with the *caret* package using repeated k-fold cross validation (5 repeats of 10-fold CV). An advantage of this resampling technique is that all observations are used for both training and validation. K = 10 folds is often used but there is no formal rule. K-fold cross validation has low bias but generally has high variance compared to other methods. Repeating k fold cross validation can be used to efficiently increase the precision of the estimates while still maintaining a small bias (Kuhn 70).

The caret package allows tuning of the hyperparameters of the models. To evaluate the models a suitable performance metrics needs to be defined. Mostly predictive accuracy is used but might not be appropriate when the data is imbalanced because a simple default strategy of guessing the majority class would give a high predictive accuracy without considering the minority class. For imbalanced datasets the binary classification measures Area Under the ROC Curve (AUC) and Cohen's Kappa are recommended. AUC summarizes the plot of true positive rate against false positive rate (the ROC curve) in a single value (Chawla 855). Kappa is the percentage of correctly classified instances out of all instances normalized at the baseline of random chance on the dataset (Brownlee). All models were trained and validated twice, once with AUC and once with Kappa as performance measure. The models trained with AUC scored higher on the test set in absolute savings (see Validation savings). Thus, AUC was selected for all models.

Feature Engineering

The raw data contains typical raw credit card fraud detection features for each transaction such as amount, date and time, merchant type (e.g. gas station), entry mode, among others (as stated above). Just with those attributes, fraud may be identified at the transactional level. However, a single transaction is not enough to detect a fraudulent transaction since it leaves behind the customer spending behavior. In order to fulfill this problem, Whitrow et al. propose to perform transaction aggregation (31-51).

Creating the aggregated features consists in grouping the transactions made during the last given number of hours by card number ("tokenized_pan"), followed by calculating the number of

transactions ("cnt_" feature) and the total amount spent on those transactions ("sum_" feature) for every transaction within the time window. We processed those new attributes for time windows of 1 day, 2 days, 1 week and 30 days, respectively, resulting in 8 new features for the model. When selecting the transactions to calculate the new features, we took two assumptions. First the own transaction is not considered as past behavior is modeled. Second the transactions must be non-fraudulent as normal customer behavior is modeled. Let's exemplify these new features for 1-day and 7-day time windows. Table 2 summarizes four transactions associated with two different card numbers (identified as 132 and 49). The first two columns (tokenized_pan and datetime) are used to compute the number of transactions within one day (cnt_1d) and within seven days (cnt_7d) time windows. The *amount* combined with cnt_1d and cnt_7d is used to compute the sum_1d and sum_7d attributes. In this example the first row has a value of 0 in every aggregated feature, because it is the first transaction and consequently there're no other transactions in its time window. The same is with the transaction of card #49 since it is the only one for that card in this example. However, in the second and third transaction of card #132 it can be seen how the quantity of transactions for the same card within the defined number of days (cnt_1d and cnt_7d) and the sum of the amount of these transactions (sum_1d and sum_7d) are calculated.

tokenized_pan	datetime	amount	cnt_1d	sum_1d	cnt_7d	sum_7d
132	2016-05-21 23:55	141,99	-	-	-	-
132	2016-05-24 09:41	6,00	-	-	1	141,99
49	2016-05-24 09:57	38,20	-	-	-	-
132	2016-05-25 08:23	6,00	1	6,00	2	147,99

Table 2: Illustrative calculation of aggegated features cnt_1d, sum_1d, cnt_7d and sum_7d

In addition to the previously mentioned features, we added two more features, indicating the fraud indexes by Issuer Bank ("id_issuer") and by Merchant ("id_merchant"), respectively. The "frd_by_Id_issuer" feature is the ratio of the number of frauds for each bank and overall frauds and "frd_by_id_merchant" is the ratio of the number of frauds for each merchant and overall frauds. The correlation matrix of all engineered features and "is_fraud" reveals a high positive correlation (0.73) between "is_fraud" and "frd_by_id_merchant" (Figure 3). This valuable feature indicates that certain merchants are associated with fraudulent activity.

Figure 3: Correlation matrix of all engineered features

Validation

The models performance is evaluated using the standard binary classification measures Area Under the Curve (AUC) and Kappa. In addition, a custom cost-based metric (savings) is used for evaluation. AUC and Kappa may not be the right evaluation criteria when evaluating fraud detection models because they implicitly assume that misclassification errors carry the same cost as the correct classified transactions. In practice, wrongly predicting a fraudulent transaction as legitimate usually carries a considerably higher financial cost than the opposite case (Bahnsen 136-137). The goal of companies, when it comes to fraud detection, is to take a decision to minimize the losses. Using a cost matrix that defines the cost for both types of misclassification error, a savings metric can be computed as the difference between the cost of using no algorithm (sum of the amounts of fraudulent transactions) and the associated cost of the predictions (Figure 4). In this experiment the highest savings of USD 561712.9 are achieved if all frauds are detected on the test set. The lowest savings is USD 0. This is when no model is used at all.

	Predicted Negative	Predicted Positive
Actual Negative	*USD 0*	*USD 20*
Actual Positive	*Total transaction amount*	*USD 20*

Figure 4: Cost matrix

Results and validation of analysis

Method	Package	Parameter	Selection	AUC	Kappa	Savings in USD
Logistic Regression	stats	-	-	0.7812106	0.665	412802.6
Decision Tree	rpart	cp	0.0002	0.8223049	0.711	447750.5
Random Forest	randomForest	mtry	13			
XGBoost		nrounds max_depth eta gamma colsample_bytree min_child_weight subsample				

Note: cross validation was repeated k-fold (5 repeats of 10-fold CV for all models) using the caret package

Table 3: Hyperparameters and validation results of models

Conclusion

References

Bahnsen, Alejandro Correa, et al. "Feature engineering strategies for credit card fraud detection."
Expert Systems with Applications, vol. 51, 2016, pp. 134–142.

Brownlee, Jason. "Machine Learning Evaluation Metrics in R." *machinelearningmastery,* 29 Feb.
2016. *machinelearningmastery.com/machine-learning-evaluation-metrics-in-r/.* Accessed 9
Dec. 2017.

Chawla, Nitesh V. "Data Mining for Imbalanced Datasets: An Overview." *Maimon O., Rokach L. (eds)
Data Mining and Knowledge Discovery Handbook.* Springer, Boston, MA, 2005.

Gorman, Ben. "A Kaggle Master Explains Gradient Boosting". *blog.kaggle,* No Free Hunch, 23 Jan.
2017.*http://blog.kaggle.com/2017/01/23/a-kaggle-master-explains-gradient-boosting/.*
Accessed 9 Dec. 2017.

Krivko ,M. "A hybrid model for plastic card fraud detection systems." *Expert Systems with
Applications*, vol. 37, 2010, pp. 6070–6076.

Kuhn, Max, and Johnson, Kjell. "Applied Predictive Modeling". Springer, New York, 2013

Maes, Sam et al. "Credit Card Fraud Detection Using Bayesian and Neural Networks". *Proceedings of
NF*, 2002.

Mahmoudi, Nader, et al. "Detecting credit card fraud by Modified Fisher Discriminant Analysis."
Expert Systems with Applications, vol. 42, 2015, pp. 2510–2516.

Prati, R.C. et al. "Class imbalance revisited: a new experimental setup to assess the performance of
treatments methods." *Knowledge and Information Systems*, vol. 45, 2015, pp. 247–270.

"The Nilson Report." David Robertson, 17 Oct. 2016,
www.nilsonreport.com/upload/content_promo/The_Nilson_Report_10-17-2016.pdf.
Accessed 02 Dec. 2017.

Visa, Sofia. "Issues in mining imbalanced data sets – a review paper." Proc. 16th Midwest Artificial Intelligence and Cognitive Science Conference, 2005, pp. 67–73.

Withrow, C. "Transaction aggregation as a strategy for credit card fraud detection". Data Mining and Knowledge Discovery, vol. 18, 2009, pp. 30-55.

Appendix

– Data visualizations, tables, etc. which support the work, but are not of primary importance

– List of data transformations, missing value imputations, outlier treatment, etc.

– List of any important assumptions not otherwise included

– Important code excerpts or algorithms used / developed if any.